

Linear_Ridge_Regression

May 10, 2019

```
In [1]: import pandas as pd                                # for working with data in Python
import numpy as np
import matplotlib.pyplot as plt                          # for visualization
from sklearn.model_selection import train_test_split
from sklearn.metrics import mean_squared_error
from sklearn import linear_model

# use Pandas to read in csv files. The pd.read_csv() method creates a DataFrame from a
train = pd.read_csv('train.csv')
test = pd.read_csv('test.csv')

print("1 \n")

# check out the size of the data
print("Train data shape:", train.shape)
print("Test data shape:", test.shape)

print("2 \n")

# look at a few rows using the DataFrame.head() method
# train.head()
print(train.head())
```

1

Train data shape: (1460, 81)

Test data shape: (1459, 80)

2

	Id	MSSubClass	MSZoning	LotFrontage	LotArea	Street	Alley	LotShape	\
0	1	60	RL	65.0	8450	Pave	NaN	Reg	
1	2	20	RL	80.0	9600	Pave	NaN	Reg	
2	3	60	RL	68.0	11250	Pave	NaN	IR1	
3	4	70	RL	60.0	9550	Pave	NaN	IR1	
4	5	60	RL	84.0	14260	Pave	NaN	IR1	

	LandContour	Utilities	...	PoolArea	PoolQC	Fence	MiscFeature	MiscVal	MoSold	\
0	Lvl	AllPub	...	0	NaN	NaN	NaN	0	2	
1	Lvl	AllPub	...	0	NaN	NaN	NaN	0	5	
2	Lvl	AllPub	...	0	NaN	NaN	NaN	0	9	
3	Lvl	AllPub	...	0	NaN	NaN	NaN	0	2	
4	Lvl	AllPub	...	0	NaN	NaN	NaN	0	12	

	YrSold	SaleType	SaleCondition	SalePrice
0	2008	WD	Normal	208500
1	2007	WD	Normal	181500
2	2008	WD	Normal	223500
3	2006	WD	Abnorml	140000
4	2008	WD	Normal	250000

[5 rows x 81 columns]

```
In [3]: plt.style.use(style='ggplot')
plt.rcParams['figure.figsize'] = (10, 6)
```

```
#####
# 2. Explore the data and engineer Features      ###
#####
```

```
print("3 \n")
```

3

```
In [4]: # to get more information like count, mean, std, min, max etc
# train.SalePrice.describe()
print (train.SalePrice.describe())
```

```
print("4 \n")
```

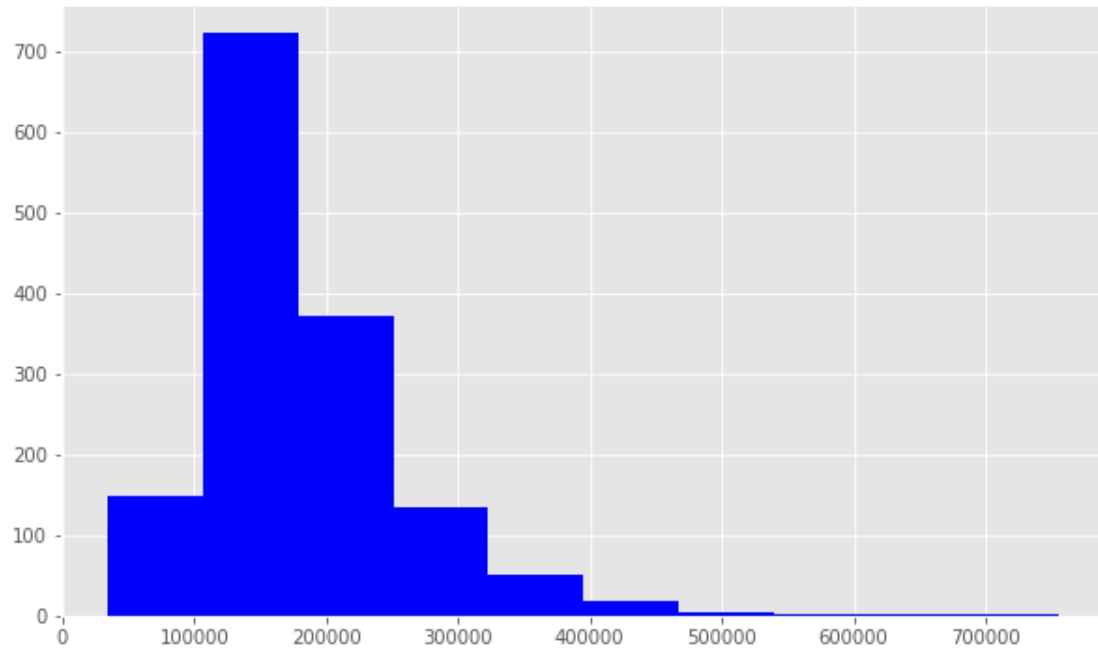
```
# to plot a histogram of SalePrice
print ("Skew is:", train.SalePrice.skew())
plt.hist(train.SalePrice, color='blue')
plt.show()
```

```
print("5 \n")
```

count	1460.000000
mean	180921.195890
std	79442.502883
min	34900.000000

```
25%      129975.000000
50%      163000.000000
75%      214000.000000
max       755000.000000
Name: SalePrice, dtype: float64
4
```

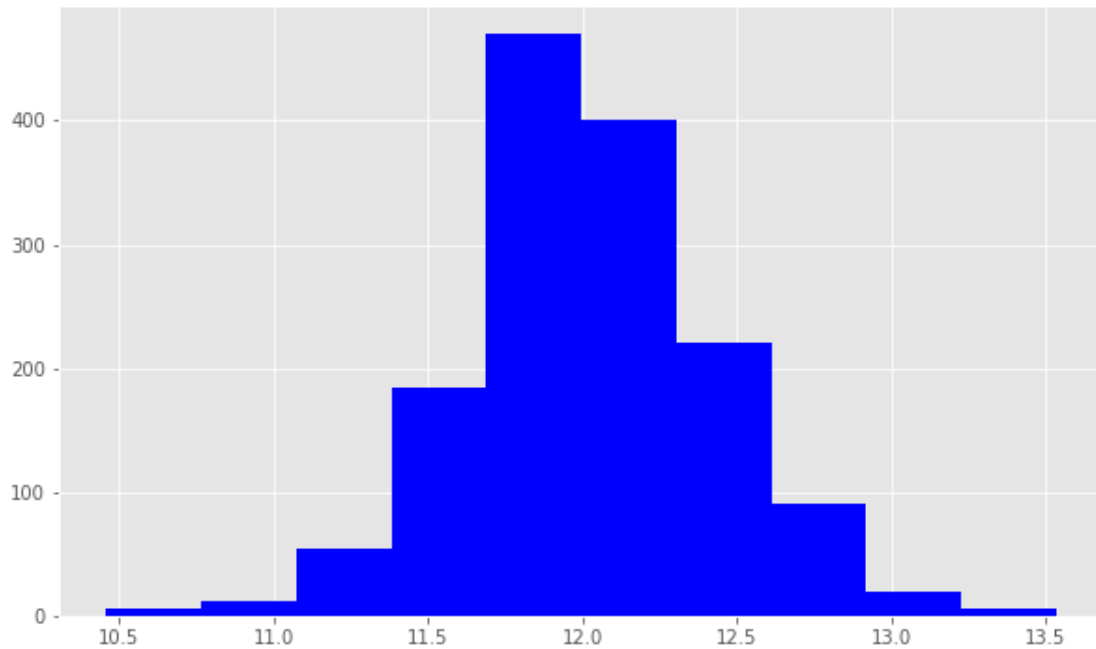
Skew is: 1.8828757597682129



5

```
In [5]: # use np.log() to transform train.SalePric and calculate the skewness a second time, a
        target = np.log(train.SalePrice)
        print ("\n Skew is:", target.skew())
        plt.hist(target, color='blue')
        plt.show()
```

Skew is: 0.12133506220520406



```
In [6]: # return a subset of columns matching the specified data types
numeric_features = train.select_dtypes(include=[np.number])
# numeric_features.dtypes
print(numeric_features.dtypes)
```

Id	int64
MSSubClass	int64
LotFrontage	float64
LotArea	int64
OverallQual	int64
OverallCond	int64
YearBuilt	int64
YearRemodAdd	int64
MasVnrArea	float64
BsmtFinSF1	int64
BsmtFinSF2	int64
BsmtUnfSF	int64
TotalBsmtSF	int64
1stFlrSF	int64
2ndFlrSF	int64
LowQualFinSF	int64
GrLivArea	int64
BsmtFullBath	int64
BsmtHalfBath	int64
FullBath	int64
HalfBath	int64

```

BedroomAbvGr      int64
KitchenAbvGr      int64
TotRmsAbvGrd      int64
Fireplaces        int64
GarageYrBltd      float64
GarageCars        int64
GarageArea        int64
WoodDeckSF        int64
OpenPorchSF       int64
EnclosedPorch     int64
3SsnPorch         int64
ScreenPorch       int64
PoolArea          int64
MiscVal           int64
MoSold            int64
YrSold            int64
SalePrice         int64
dtype: object

```

```
In [7]: corr = numeric_features.corr()
```

```

# The first five features are the most positively correlated with SalePrice, while the
print (corr['SalePrice'].sort_values(ascending=False)[:5], '\n')
print (corr['SalePrice'].sort_values(ascending=False)[-5:])

```

```

SalePrice      1.000000
OverallQual    0.790982
GrLivArea      0.708624
GarageCars     0.640409
GarageArea     0.623431
Name: SalePrice, dtype: float64

```

```

YrSold         -0.028923
OverallCond    -0.077856
MSSubClass     -0.084284
EnclosedPorch  -0.128578
KitchenAbvGr   -0.135907
Name: SalePrice, dtype: float64

```

```
In [8]: print(train.OverallQual.unique())
```

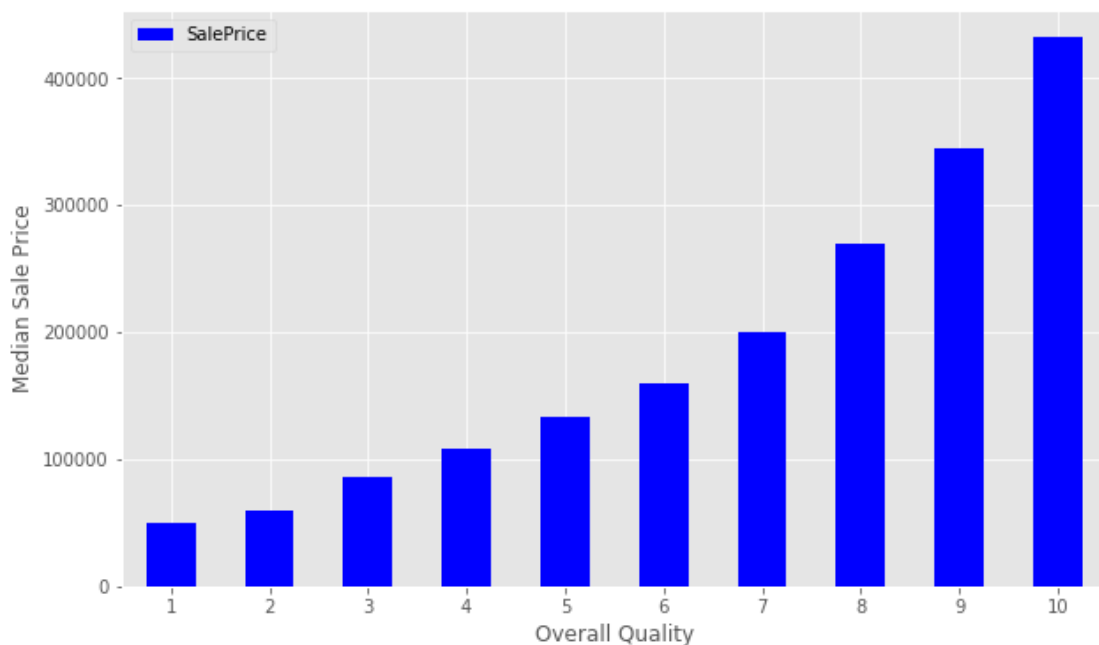
```

"""
print("9 \n")
"""
#investigate the relationship between OverallQual and SalePrice.
#We set index='OverallQual' and values='SalePrice'. We chose to look at the median here
quality_pivot = train.pivot_table(index='OverallQual', values='SalePrice', aggfunc=np.median)
print(quality_pivot)

```

```
[ 7  6  8  5  9  4 10  3  1  2]
      SalePrice
OverallQual
1           50150
2           60000
3           86250
4          108000
5          133000
6          160000
7          200141
8          269750
9          345000
10         432390
```

```
In [11]: #visualize this pivot table more easily, we can create a bar plot
         #Notice that the median sales price strictly increases as Overall Quality increases.
         quality_pivot.plot(kind='bar', color='blue')
         plt.xlabel('Overall Quality')
         plt.ylabel('Median Sale Price')
         plt.xticks(rotation=0)
         plt.show()
```



```
In [12]: print("11 \n")
         """
         #to generate some scatter plots and visualize the relationship between the Ground Liv
```

```

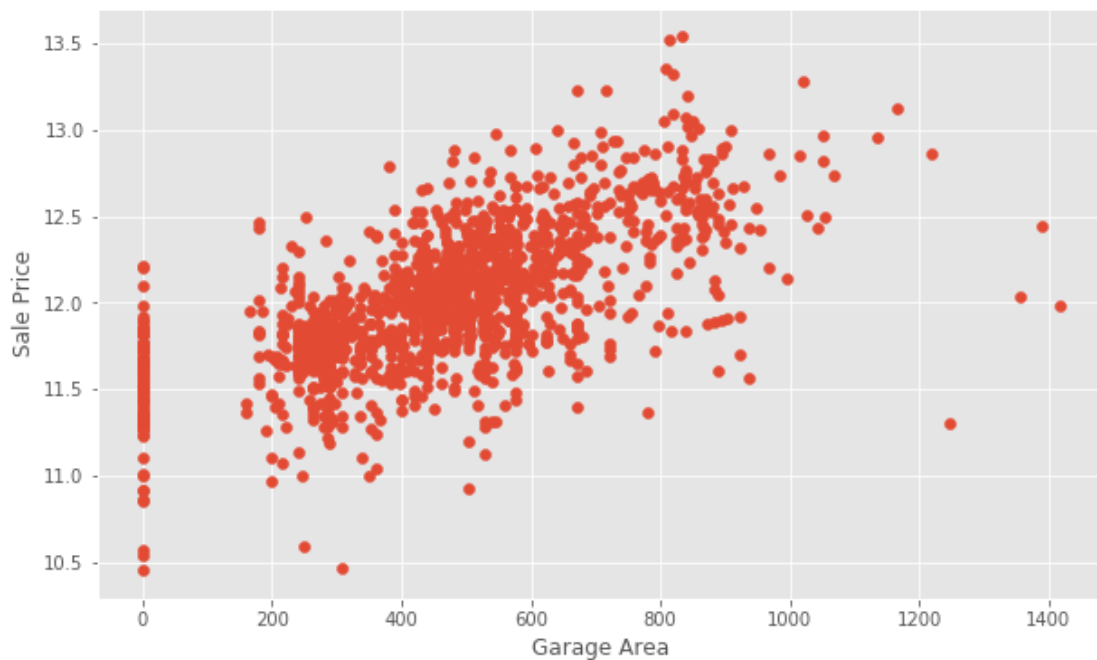
plt.scatter(x=train['GrLivArea'], y=target)
plt.ylabel('Sale Price')
plt.xlabel('Above grade (ground) living area square feet')
plt.show()
"""
print("12 \n")

# do the same for GarageArea.
plt.scatter(x=train['GarageArea'], y=target)
plt.ylabel('Sale Price')
plt.xlabel('Garage Area')
plt.show()

```

11

12



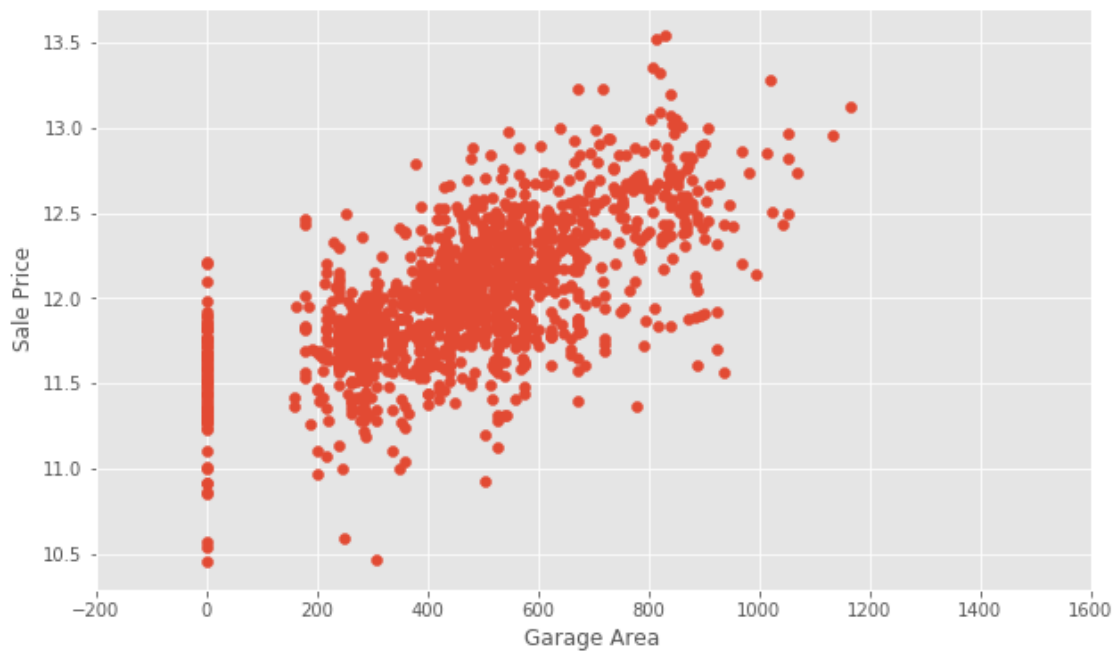
```

In [13]: # create a new dataframe with some outliers removed
train = train[train['GarageArea'] < 1200]

# display the previous graph again without outliers
plt.scatter(x=train['GarageArea'], y=np.log(train.SalePrice))
plt.xlim(-200,1600)      # This forces the same scale as before
plt.ylabel('Sale Price')

```

```
plt.xlabel('Garage Area')
plt.show()
```



```
In [14]: # create a DataFrame to view the top null columns and return the counts of the null v
nulls = pd.DataFrame(train.isnull().sum().sort_values(ascending=False)[:25])
nulls.columns = ['Null Count']
nulls.index.name = 'Feature'
#nulls
print(nulls)
```

Feature	Null Count
PoolQC	1449
MiscFeature	1402
Alley	1364
Fence	1174
FireplaceQu	689
LotFrontage	258
GarageCond	81
GarageType	81
GarageYrBlt	81
GarageFinish	81
GarageQual	81
BsmtExposure	38
BsmtFinType2	38
BsmtFinType1	37

BsmtCond	37
BsmtQual	37
MasVnrArea	8
MasVnrType	8
Electrical	1
Utilities	0
YearRemodAdd	0
MSSubClass	0
Foundation	0
ExterCond	0
ExterQual	0

```
In [15]: print("15 \n")
        """
        #to return a list of the unique values
        print ("Unique values are:", train.MiscFeature.unique())
        """

        #####
        #   Wrangling the non-numeric Features   ##
        #####

        print("16 \n")

        # consider the non-numeric features and display details of columns
        categoricals = train.select_dtypes(exclude=[np.number])
        #categoricals.describe()
        print(categoricals.describe())
```

15

16

	MSZoning	Street	Alley	LotShape	LandContour	Utilities	LotConfig	\
count	1455	1455	91	1455	1455	1455	1455	
unique	5	2	2	4	4	2	5	
top	RL	Pave	Grvl	Reg	Lvl	AllPub	Inside	
freq	1147	1450	50	921	1309	1454	1048	

	LandSlope	Neighborhood	Condition1	...	GarageType	GarageFinish	\
count	1455		1455	1455	...	1374	1374
unique	3		25	9	...	6	3
top	Gtl		NAmes	Norm	...	Attchd	Unf
freq	1378		225	1257	...	867	605

	GarageQual	GarageCond	PavedDrive	PoolQC	Fence	MiscFeature	SaleType	\
count	1374	1374	1455	6	281	53	1455	

unique	5	5	3	3	4	4	9
top	TA	TA	Y	Gd	MnPrv	Shed	WD
freq	1306	1321	1335	2	157	48	1266

SaleCondition	
count	1455
unique	6
top	Normal
freq	1196

[4 rows x 43 columns]

```
In [16]: #####
#   Transforming and engineering features           ##
#####

print("17 \n")

# When transforming features, it's important to remember that any transformations tha
# fitting the model must be applied to the test data.

#Eg:
print ("Original: \n")
print (train.Street.value_counts(), "\n")
```

17

Original:

Pave	1450
Grvl	5

Name: Street, dtype: int64

```
In [17]: print("18 \n")

# our model needs numerical data, so we will use one-hot encoding to transform the da
# create a new column called enc_street. The pd.get_dummies() method will handle this
train['enc_street'] = pd.get_dummies(train.Street, drop_first=True)
test['enc_street'] = pd.get_dummies(test.Street, drop_first=True)

print ('Encoded: \n')
print (train.enc_street.value_counts()) # Pave and Grvl values converted into 1 and

print("19 \n")
```

```

# look at SaleCondition by constructing and plotting a pivot table, as we did above f
condition_pivot = train.pivot_table(index='SaleCondition', values='SalePrice', aggfun
condition_pivot.plot(kind='bar', color='blue')
plt.xlabel('Sale Condition')
plt.ylabel('Median Sale Price')
plt.xticks(rotation=0)
plt.show()

```

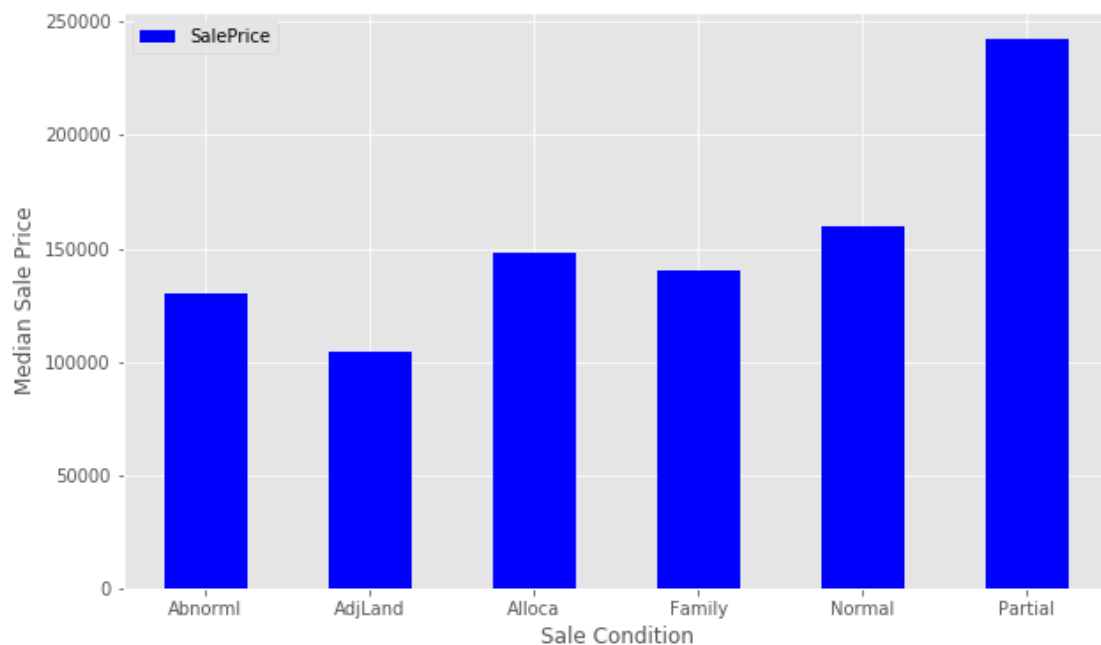
18

Encoded:

```

1    1450
0      5
Name: enc_street, dtype: int64
19

```



```

In [18]: # encode this SaleCondition as a new feature by using a similar method that we used f
def encode(x): return 1 if x == 'Partial' else 0
train['enc_condition'] = train.SaleCondition.apply(encode)
test['enc_condition'] = test.SaleCondition.apply(encode)

print("20 \n")

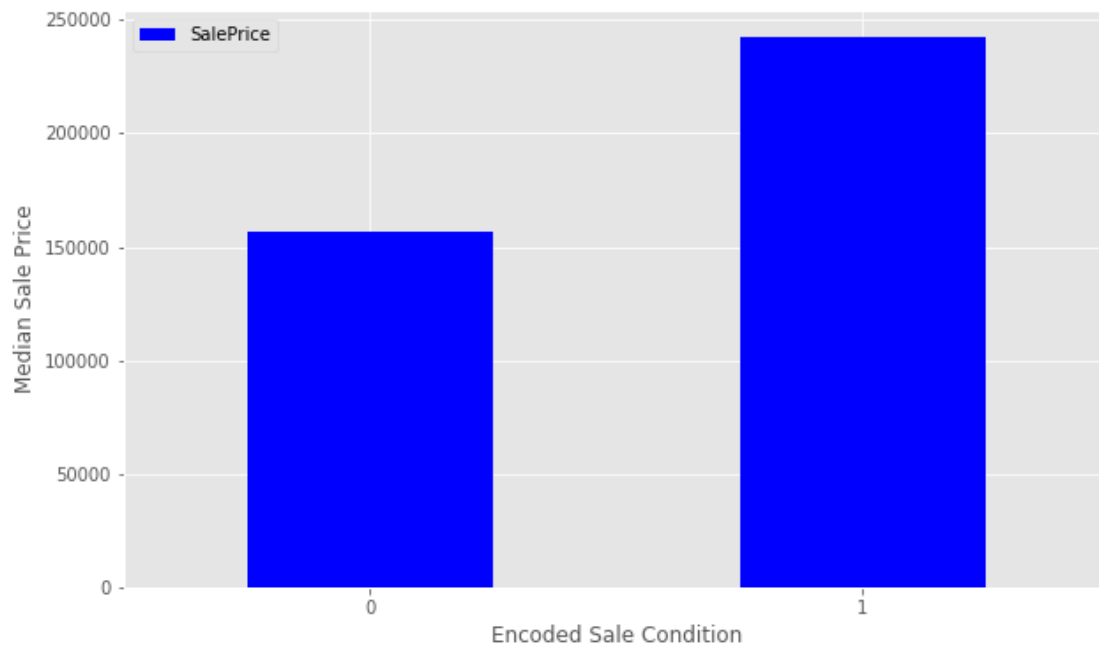
```

```

# explore this newly modified feature as a plot.
condition_pivot = train.pivot_table(index='enc_condition', values='SalePrice', aggfun
condition_pivot.plot(kind='bar', color='blue')
plt.xlabel('Encoded Sale Condition')
plt.ylabel('Median Sale Price')
plt.xticks(rotation=0)
plt.show()

```

20



```

In [19]: #####
#   Dealing with missing values
#   We'll fill the missing values with an average value and then assign the results t
#   This is a method of interpolation
#####
data = train.select_dtypes(include=[np.number]).interpolate().dropna()

print("21 \n")
# Check if the all of the columns have 0 null values.
# sum(data.isnull().sum() != 0)
print(sum(data.isnull().sum() != 0))

print("22 \n")

```

21

0

22

```
In [20]: #####
# 3. Build a linear model                                     ##
#####

# separate the features and the target variable for modeling.
# We will assign the features to X and the target variable(Sale Price)to y.

y = np.log(train.SalePrice)
X = data.drop(['SalePrice', 'Id'], axis=1)
# exclude ID from features since Id is just an index with no relationship to SalePrice

#===== partition the data =====
# Partitioning the data in this way allows us to evaluate how our model might perform
# If we train the model on all of the test data, it will be difficult to tell if overfitting
#=====
# also state how many percentage from train data set, we want to take as test data set
# In this example, about 33% of the data is devoted to the hold-out set.
X_train, X_test, y_train, y_test = train_test_split(X, y, random_state=42, test_size=0.33)

In [21]: #===== Begin modelling =====#
# Linear Regression Model                                     #
#=====#

# ---- first create a Linear Regression model.
# First, we instantiate the model.
lr = linear_model.LinearRegression()

# ---- fit the model / Model fitting
# lr.fit() method will fit the linear regression on the features and target variable
model = lr.fit(X_train, y_train)

print("23 \n")
```

23

```
In [22]: # ---- Evaluate the performance and visualize results
# r-squared value is a measure of how close the data are to the fitted regression line
# a higher r-squared value means a better fit(very close to value 1)
print("R^2 is: \n", model.score(X_test, y_test))
```

```

# use the model we have built to make predictions on the test data set.
predictions = model.predict(X_test)

print("24 \n")

```

R² is:
0.8882477709262542
24

```

In [23]: print('RMSE is: \n', mean_squared_error(y_test, predictions))

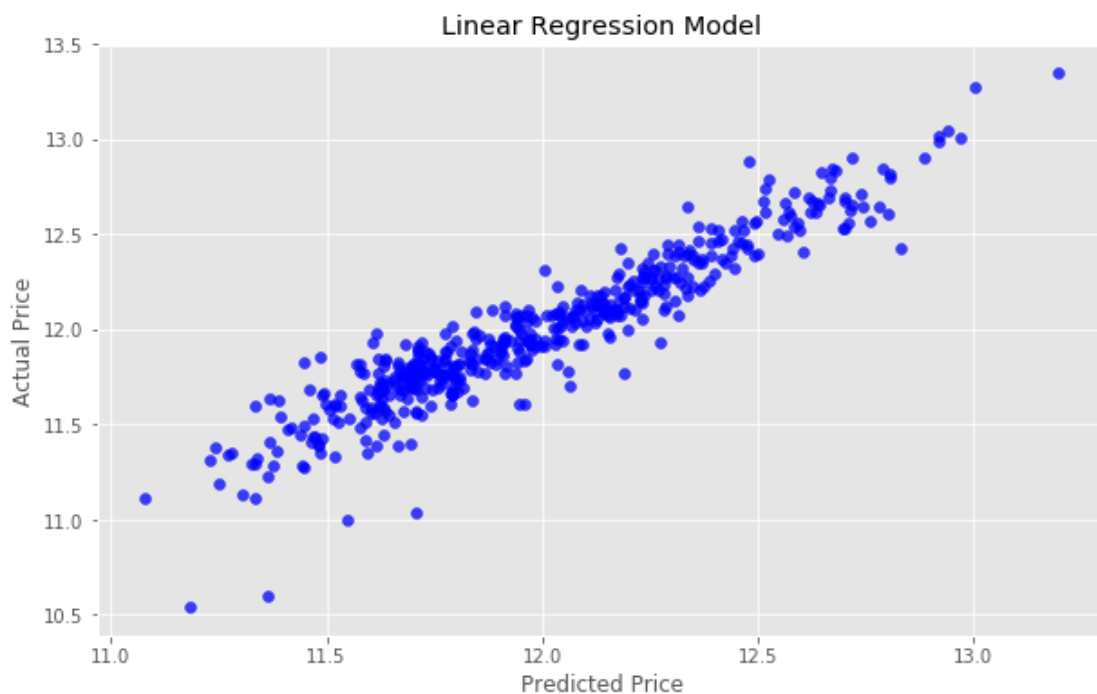
```

```

print("25 \n")
# view this relationship between predictions and actual_values graphically with a scatter plot
actual_values = y_test
plt.scatter(predictions, actual_values, alpha=.75,
            color='b') # alpha helps to show overlapping data
plt.xlabel('Predicted Price')
plt.ylabel('Actual Price')
plt.title('Linear Regression Model')
plt.show()

```

RMSE is:
0.017841794519567734
25



```

In [24]: #===== improve the model =====
# try using Ridge Regularization to decrease the influence of less important feature.
#=====

print("26 \n")
# experiment by looping through a few different values of alpha, and see how this cha

for i in range (-2, 3):
    alpha = 10**i
    rm = linear_model.Ridge(alpha=alpha)
    ridge_model = rm.fit(X_train, y_train)
    preds_ridge = ridge_model.predict(X_test)

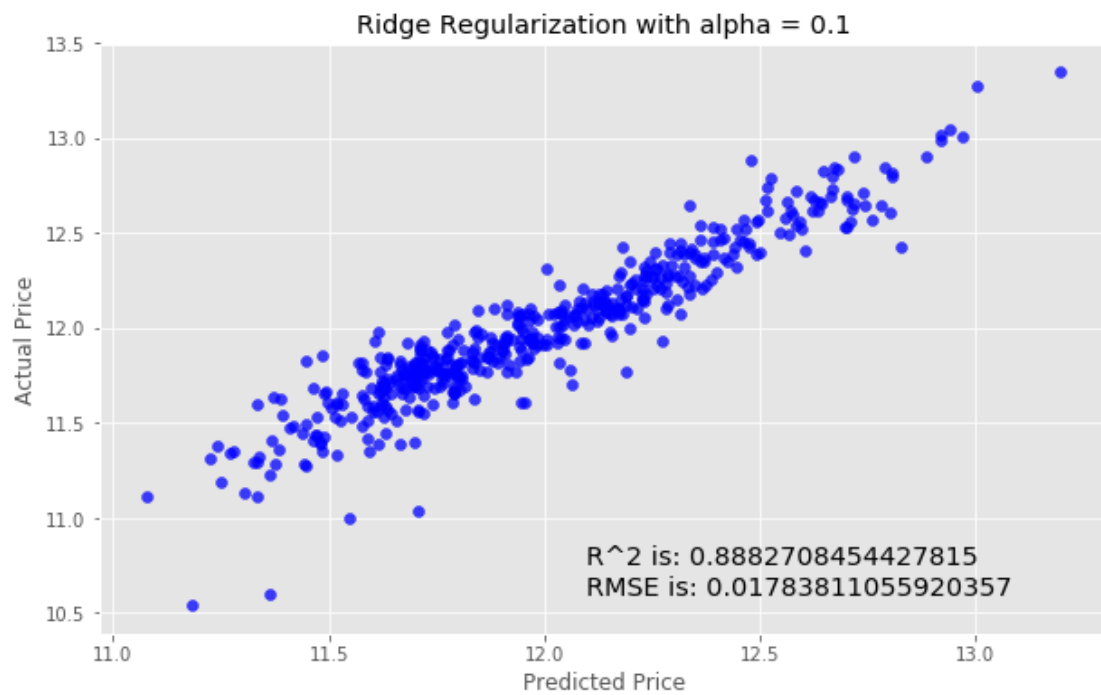
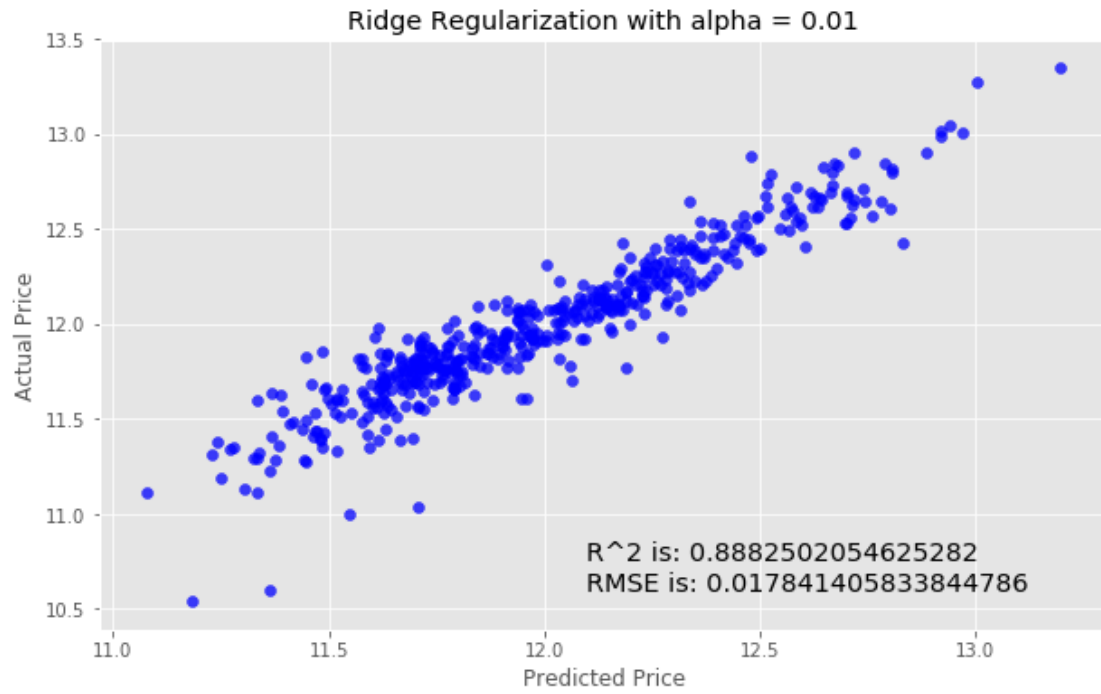
    plt.scatter(preds_ridge, actual_values, alpha=.75, color='b')
    plt.xlabel('Predicted Price')
    plt.ylabel('Actual Price')
    plt.title('Ridge Regularization with alpha = {}'.format(alpha))
    overlay = 'R^2 is: {}\nRMSE is: {}'.format(
        ridge_model.score(X_test, y_test),
        mean_squared_error(y_test, preds_ridge))
    plt.annotate(s=overlay,xy=(12.1,10.6),size='x-large')
    plt.show()

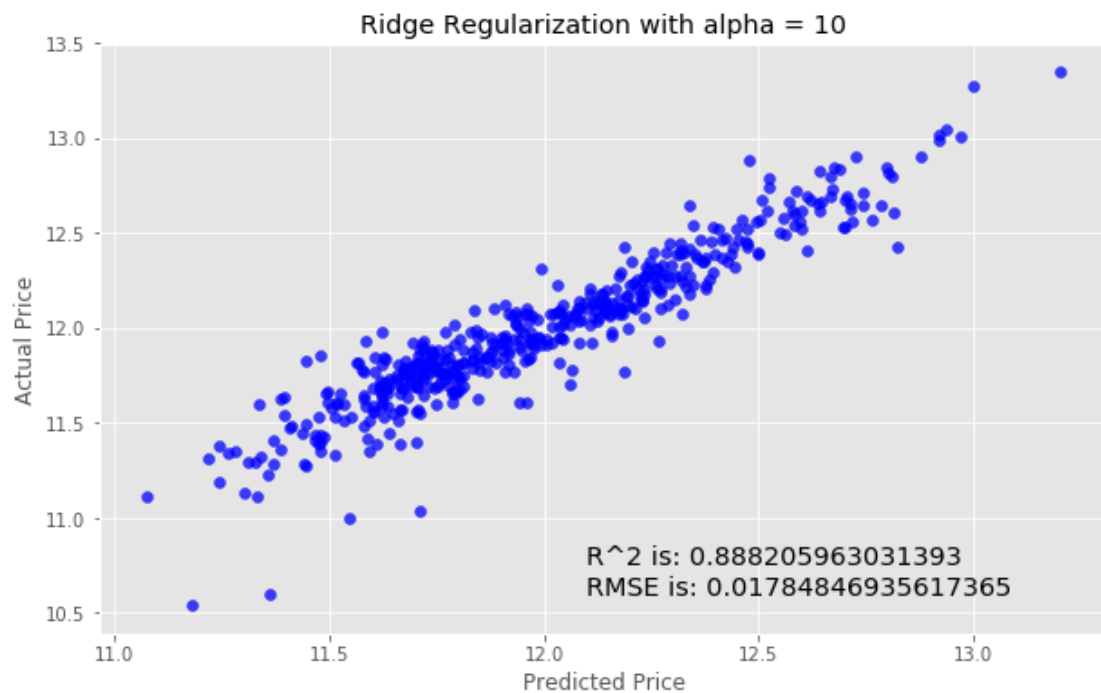
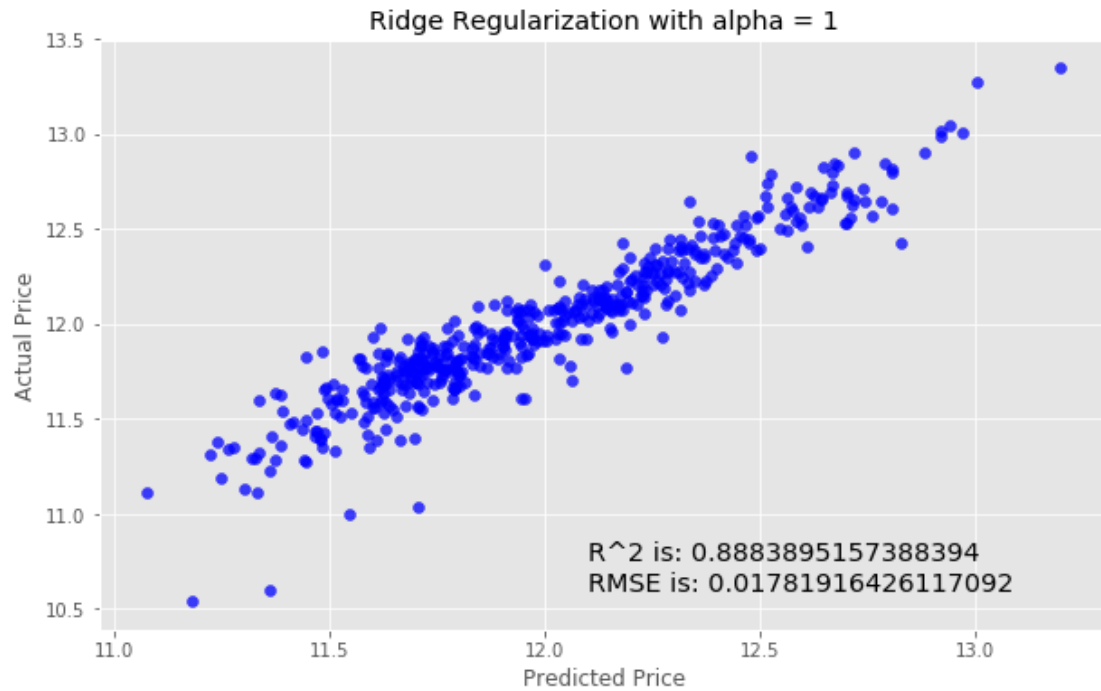
# if you examined the plots you can see these models perform almost identically to th
# In our case, adjusting the alpha did not substantially improve our model.

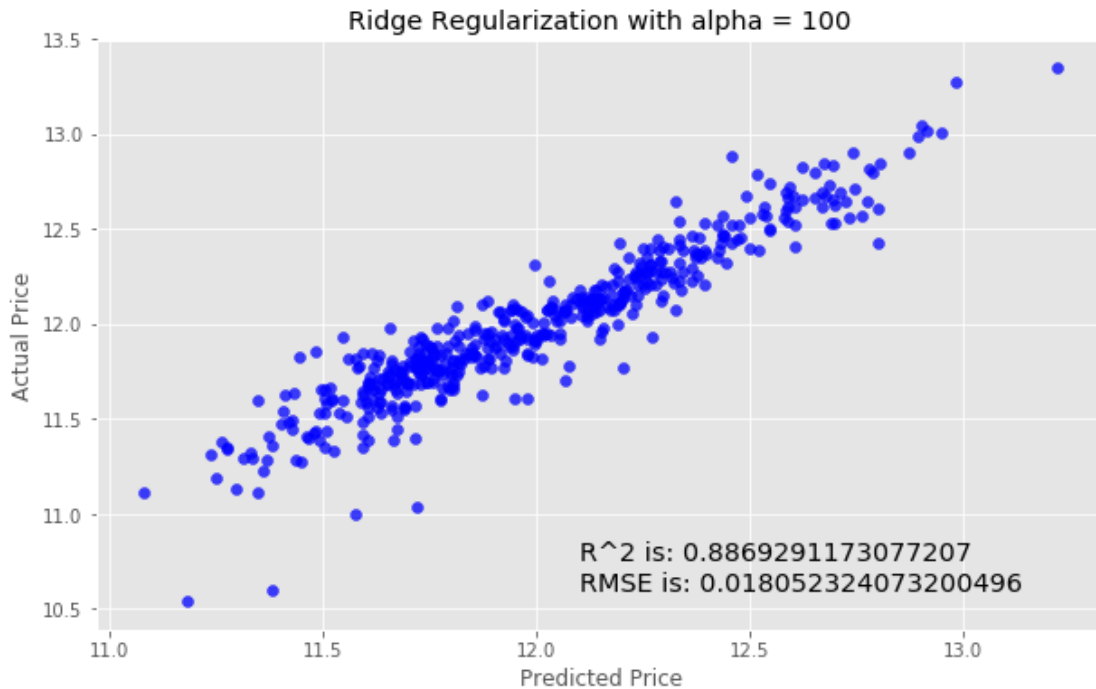
print("27 \n")
print("R^2 is: \n", model.score(X_test, y_test))

```

26







27

R^2 is:
0.8882477709262542

```
In [25]: #####
# 4. Make a submission ##
#####

# create a csv that contains the predicted SalePrice for each observation in the test
submission = pd.DataFrame()
# The first column must contain the ID from the test data.
submission['Id'] = test.Id

# select the features from the test data for the model as we did above.
feats = test.select_dtypes(
    include=[np.number]).drop(['Id'], axis=1).interpolate()

# generate predictions
predictions = model.predict(feats)

# transform the predictions to the correct form
# apply np.exp() to our predictions because we have taken the logarithm(np.log()) pr
final_predictions = np.exp(predictions)
```

```

print("28 \n")

# check the difference
print("Original predictions are: \n", predictions[:10], "\n")
print("Final predictions are: \n", final_predictions[:10])

print("29 \n")
# assign these predictions and check
submission['SalePrice'] = final_predictions
# submission.head()
print(submission.head())

# export to a .csv file as Kaggle expects.
# pass index=False because Pandas otherwise would create a new index for us.
submission.to_csv('submission1.csv', index=False)

print("\n Finish")

```

28

Original predictions are:

```
[11.76725362 11.71929504 12.07656074 12.20632678 12.11217655 12.05709882
12.16036698 12.01665734 12.17126892 11.66318882]
```

Final predictions are:

```
[128959.49172586 122920.74024359 175704.82598102 200050.83263756
182075.46986405 172318.33397533 191064.62164201 165488.55901671
193158.99133192 116214.02546462]
```

29

	Id	SalePrice
0	1461	128959.491726
1	1462	122920.740244
2	1463	175704.825981
3	1464	200050.832638
4	1465	182075.469864

Finish